How many of each cause of homicide?

First download the `Baltimore_homicides.zip` file from Courseplus and unzip it into your working directory. You can read the file `homicides.txt` with `readLines` via

```
homicides <- readLines("homicides.txt")
```

The goal of this exercise is to count the number of different types of homicides are in this dataset. In each record there is a field with the word "Cause" in it indicating the cause of death (e.g. "Cause: shooting"). Try to extract this field and count the number of instances of each cause.

The answer should look something like

```
asphyxiation  blunt force        other     shooting     stabbing      unknown
          28           77            6         1003          121           10
```

Age distribution of homicide victims

The goal of this part is to make a histogram of the age distribution of the homicide victims in the dataset. For most (but not all) records there is an indication of the age of the victim, however it is not always consisten. Extract the age of the victim using a regular expression and make a histogram of the ages. It should look something like the following:

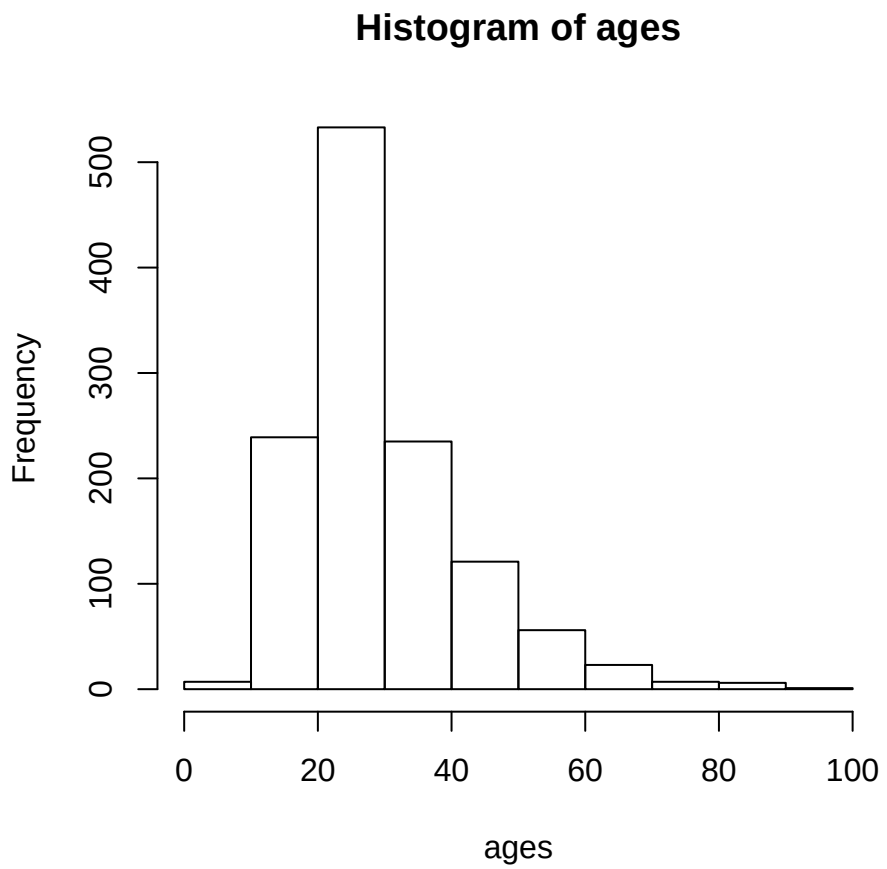

Spatial distribution of homicides

Each homicide in the dataset contains a latitude/longitude pair identifying the location (the block) where the victim was found. Make a map of the locations of the homicides to examine the spatial distribution of where homicide victims are found. You can use the **maps** package to draw a map of Baltimore.

```
library(maps)
map("county", "maryland,baltimore city")
```

The resulting map should look something like the following.