

Reproducible Research and Literate Statistical Programming (with knitr)

Roger D. Peng, PhD

Department of Biostatistics

Johns Hopkins Bloomberg School of Public Health

University of Minnesota

April 2013


Replication

- The ultimate standard for strengthening scientific evidence is replication of findings and conducting studies with independent
 - Investigators
 - Data
 - Analytical methods
 - Laboratories
 - Instruments
- Replication is particularly important in studies that can impact broad policy or regulatory decisions

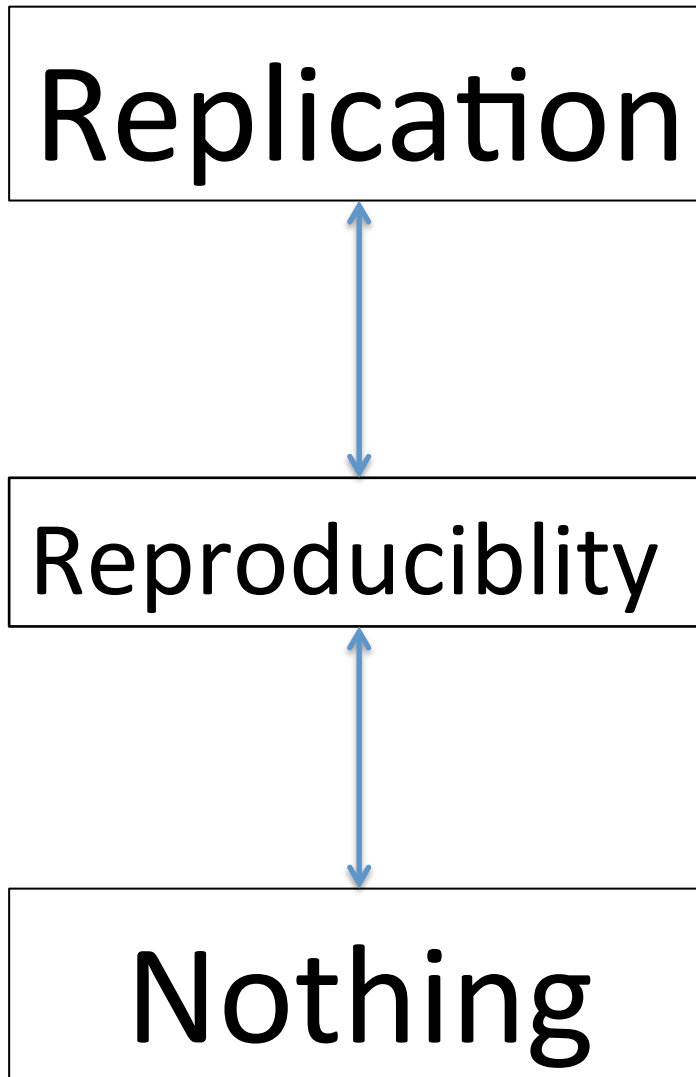
What's Wrong with Replication?

- Some studies cannot be replicated
 - No time, opportunistic
 - No money
 - Unique
- Reproducible Research: Make analytic data and code available so that others may reproduce findings

How Can We Bridge the Gap?



How Can We Bridge the Gap?



Why Do We Need Reproducible Research Now?

- New technologies increasing data collection throughput; data are more complex and extremely high dimensional
- Existing databases can be merged into new “megadatabases”
- Computing power is greatly increased, allowing more sophisticated analyses
- For every field “X” there is a field “Computational X”

Example: Reproducible Air Pollution and Health Research

- Estimating small (but important) health effects in the presence of much stronger signals
- Results inform substantial policy decisions, affect many stakeholders
 - EPA regulations can cost billions of dollars
- Complex statistical methods are needed and subjected to intense scrutiny

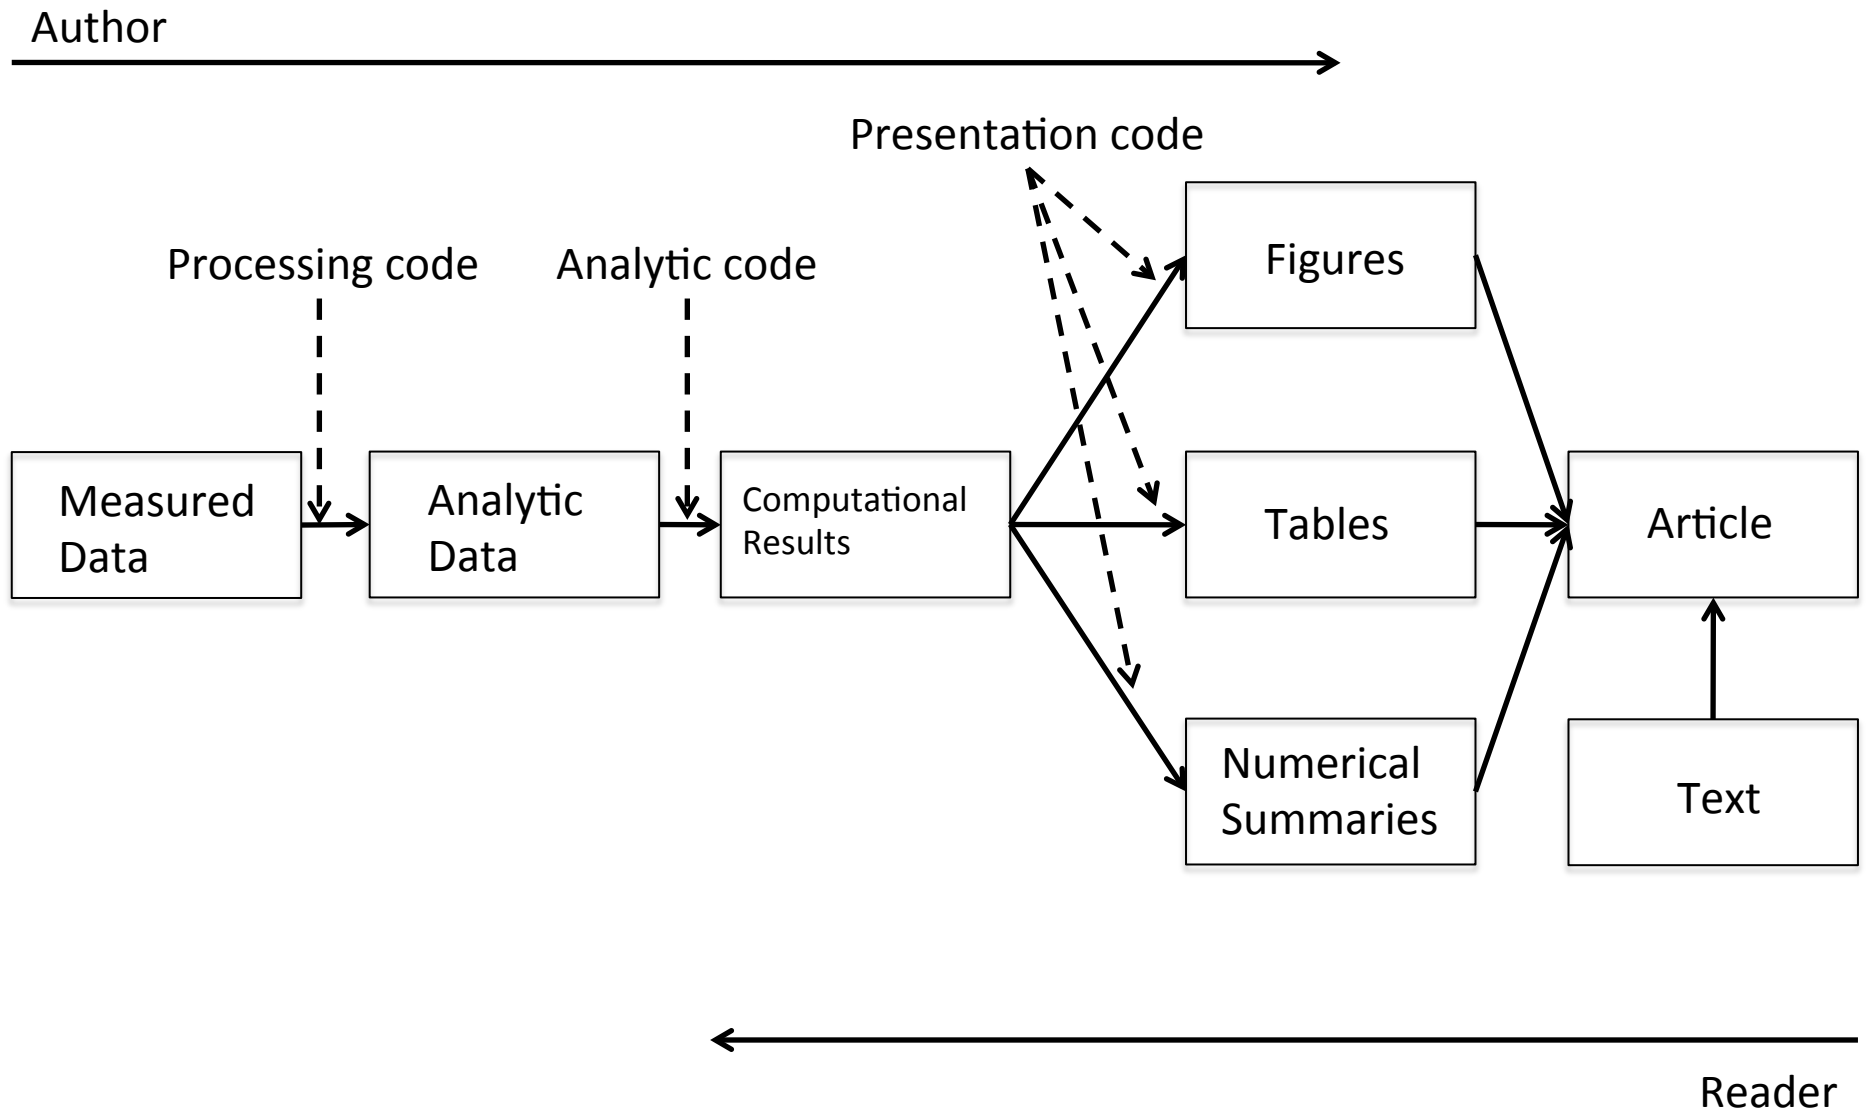
Research Pipeline



Article

Reader

Research Pipeline



Recent Developments in Reproducible Research



Data Replication & Reproducibility

PERSPECTIVE

Reproducible Research in Computational Science

Roger D. Peng

Computational science has led to exciting new developments, but the nature of the work has exposed limitations in our ability to evaluate published findings. Reproducibility has the potential to serve as a minimum standard for judging scientific claims when full independent replication of a study is not possible.

Recent Developments in Reproducible Research

Economix



Explaining the Science of Everyday Life

April 17, 2013, 3:28 pm |  19 Comments

With Debt Study's Errors Confirmed, Debate on Conclusion Goes On

By ANNIE LOWREY

The Harvard economists Carmen M. Reinhart and Kenneth S. Rogoff [have acknowledged that](#) their groundbreaking 2010 study “[Growth in a Time of Debt](#)” includes statistical errors that significantly alter its results.

Recent Developments in Reproducible Research

The Duke Saga

<http://goo.gl/hijaN>



Recent Developments in Reproducible Research

REPORT BRIEF  MARCH 2012


INSTITUTE OF MEDICINE
OF THE NATIONAL ACADEMIES

Advising the nation • Improving health

For more information visit www.iom.edu/translationalomics

Evolution of Translational Omics

Lessons Learned and the
Path Forward



<http://goo.gl/ZlIJa>

The IOM Report

In the Discovery/Test Validation stage of omics-based tests:

- **Data/metadata** used to develop test should be made publicly available
- The **computer code** and fully specified computational procedures used for development of the candidate omics-based test should be made sustainably available
- “Ideally, the computer code that is released will **encompass all of the steps of computational analysis**, including all data preprocessing steps, that have been described in this chapter. All aspects of the analysis need to be transparently reported.”

When Can Research be Reproducible?

- Analytic data are available
- Analytic code are available
- Standard means of distribution
- Documentation of code and data

Literate (Statistical) Programming

- An article is a stream of **text** and **code**
- Analysis code is divided into text and code “chunks”
- Each code chunk loads data and computes results
- Presentation code formats results (tables, figures, etc.)
- Article text explains what is going on
- Literate programs can be **weaved** to produce human-readable documents and **tangled** to produce machine-readable documents

Literate (Statistical) Programming

- Literate programming is a general concept that requires
 1. A documentation language (human readable)
 2. A programming language (machine readable)
- Sweave uses L^AT_EX and R as the documentation and programming languages
- Sweave was developed by Friedrich Leisch (member of the R Core) and is maintained by R core
- Main web site: `http://www.statistik.lmu.de/~leisch/Sweave`

Sweave Limitations

- Sweave has many limitations
- Focused primarily on LaTeX, a difficult to learn markup language used only by weirdos
- Lacks features like caching, multiple plots per chunk, mixing programming languages and many other technical items
- Not frequently updated or very actively developed

Literate (Statistical) Programming

- **knitr** is an alternative (more recent) package
- Brings together many features added on to Sweave to address limitations
- knitr uses R as the programming language (although others are allowed) and variety of documentation languages
 - LaTeX, Markdown, HTML
- knitr was developed by Yihui Xie (graduate student in statistics at Iowa State)
- See <http://yihui.name/knitr/>

What is LSP good for?

- Manuals
- Short/medium-length technical documents
- Tutorials
- Reports (esp. if generated periodically)
- Data preprocessing documents/summaries

What is LSP NOT good for?

- Long research articles
- Complex time-consuming computations (i.e. long MCMC, optimizations)
- Documents that require precise formatting
- Books (although can be done with use of other tools)

Using knitr


- Toolchain
 - Text editor
 - R
 - knitr package (from CRAN via `install.packages`)
 - An appropriate document viewer (PDF viewer, web browser)
 - pandoc (optional)
- Easiest to use RStudio where knitr (and Sweave) are already integrated

knitr

[Home](#)[Objects](#)[Options](#)[Hooks](#)[Patterns](#)[Demos](#)

knitr

Elegant, flexible and fast dynamic report generation with R



Overview

The **knitr** package was designed to be a transparent engine for dynamic report generation with R, solve some long-standing problems in Sweave, and combine features in other add-on packages into one package (**knitr** \approx Sweave + cacheSweave + pgfSweave + weaver + animation::saveLatex + R2HTML::RweaveHTML + highlight::HighlightWeaveLatex + 0.2 * brew + 0.1 * SweaveListingUtils + more).

<http://yihui.name/knitr>

Pandoc

Pandoc a universal document converter



85

[About](#)

[Installing](#)

[Demos](#)

[Try pandoc online](#)

[Examples](#)

[Documentation](#)

[User's Guide](#)

[FAQ](#)

[API documentation](#)

[Contributing](#)

[Mailing lists](#)

[Scripting](#)

[Making an ebook](#)

[Extras](#)

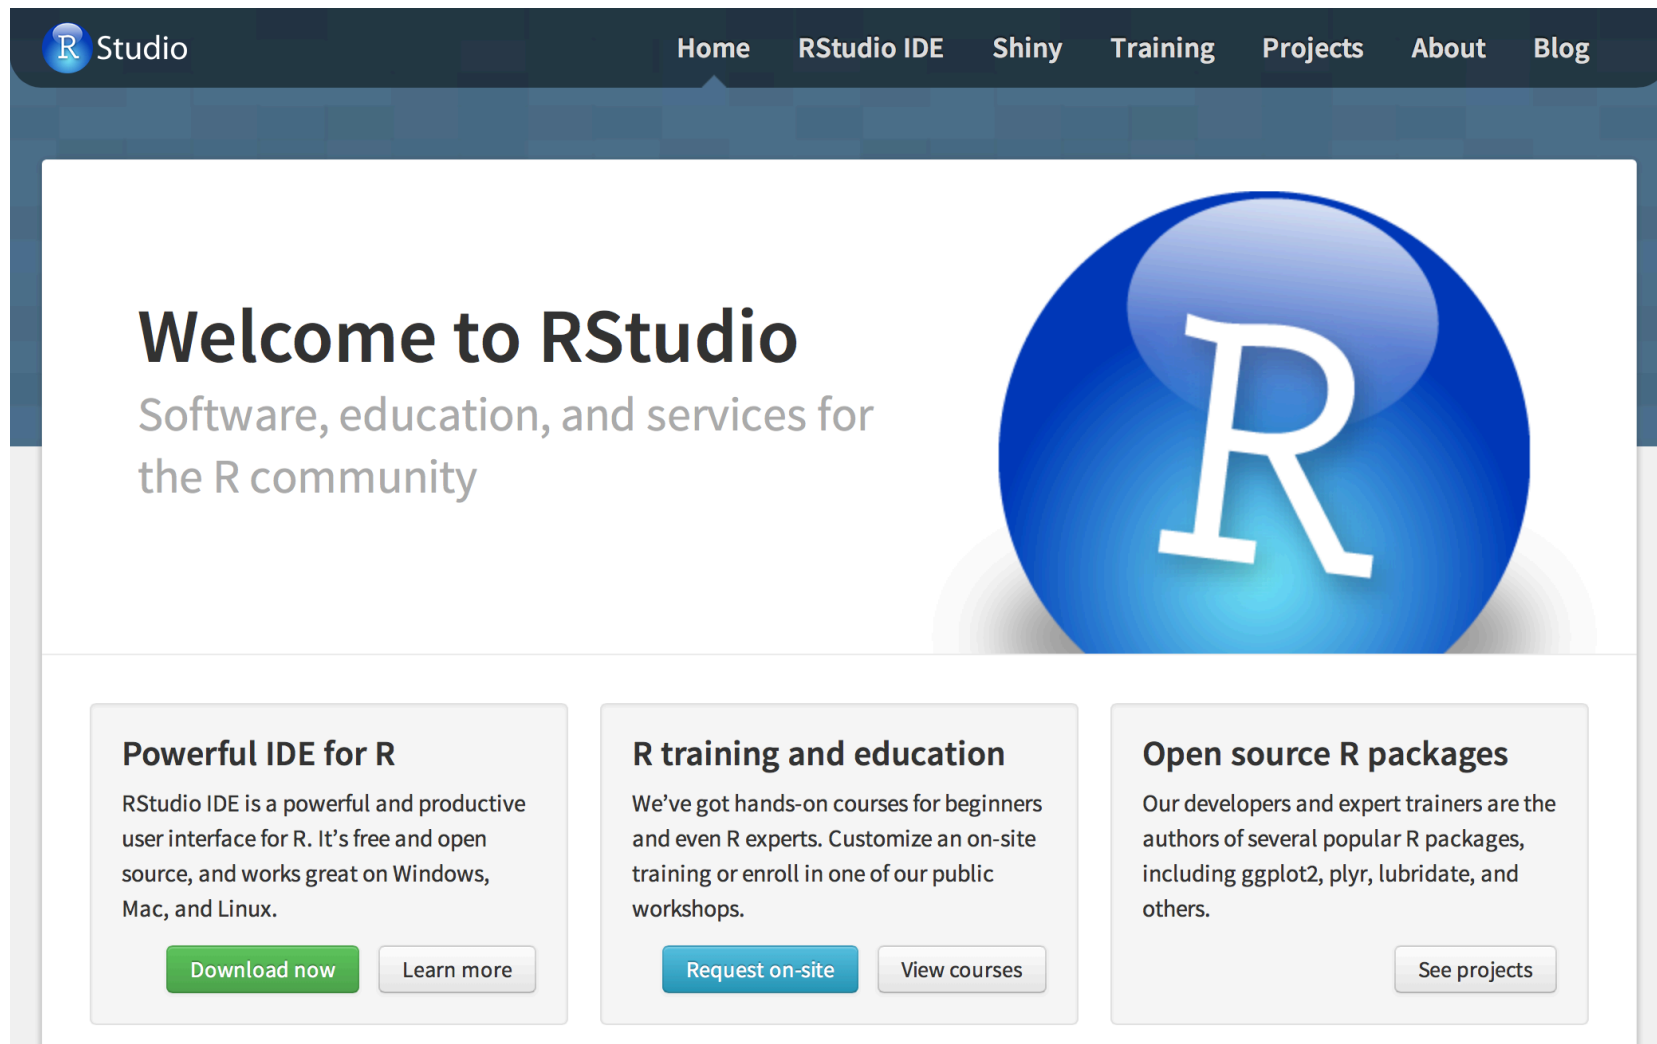
[Releases](#)

About pandoc

If you need to convert files from one markup format into another, pandoc is your swiss-army knife. Pandoc can convert documents in [markdown](#), [reStructuredText](#), [textile](#), [HTML](#), [DocBook](#), [LaTeX](#), or [MediaWiki markup](#) to

- HTML formats: XHTML, HTML5, and HTML slide shows using [Slidy](#), [Slideous](#), [S5](#), or [DZSlides](#).
- Word processor formats: Microsoft Word [docx](#), OpenOffice/LibreOffice [ODT](#), [OpenDocument XML](#)
- Ebooks: [EPUB](#) version 2 or 3, [FictionBook2](#)
- Documentation formats: [DocBook](#), [GNU TexInfo](#), [Groff man](#) pages
- TeX formats: [LaTeX](#), [ConTeXt](#), LaTeX Beamer slides
- [PDF](#) via LaTeX
- Lightweight markup formats: [Markdown](#), [reStructuredText](#), [AsciiDoc](#), [MediaWiki markup](#), Emacs [Org-Mode](#), [Textile](#)

RStudio

The image is a screenshot of the RStudio website homepage. At the top, there is a dark blue navigation bar with the RStudio logo on the left and a series of links: Home, RStudio IDE, Shiny, Training, Projects, About, and Blog. Below the navigation bar, the main content area has a white background. On the left, the text 'Welcome to RStudio' is followed by 'Software, education, and services for the R community'. To the right of this text is a large, 3D-style blue sphere with a white 'R' on it. Below the main heading, there are three light gray boxes. The first box is titled 'Powerful IDE for R' and describes the RStudio IDE as a powerful and productive user interface for R, free and open source, and works on Windows, Mac, and Linux. It has two buttons: 'Download now' (green) and 'Learn more' (gray). The second box is titled 'R training and education' and describes hands-on courses for beginners and even R experts, with options for on-site training or public workshops. It has two buttons: 'Request on-site' (blue) and 'View courses' (gray). The third box is titled 'Open source R packages' and describes the developers and expert trainers as authors of several popular R packages, including ggplot2, plyr, lubridate, and others. It has one button: 'See projects' (gray).

<http://www.rstudio.com>

Markdown

 **DARING FIREBALL**

By JOHN GRUBER

[ARCHIVE](#)
[MEMBERSHIP](#)
[PROJECTS](#)
[CONTACT](#)
[COLOPHON](#)
[RSS FEED](#)
[SPONSORSHIP](#)

Markdown

[Main](#) [Basics](#) [Syntax](#) [License](#) [Dingus](#)

DOWNLOAD

[Markdown 1.0.1](#) (18 KB) — 17 Dec 2004

INTRODUCTION

Markdown is a text-to-HTML conversion tool for web writers. Markdown allows you to write using an easy-to-read, easy-to-write plain text format, then convert it to structurally valid XHTML (or HTML).

<http://daringfireball.net/projects/markdown/>

Markdown

- A simpler version of “markup” languages
- Written in simple text format
- Has minimal formatting indicators
- Tools support converting it to many other formats (pandoc)

Markdown

- # indicates a first-level header
- ## second-level header
- ### third-level header
- Indentation can be used to create an ordered (with numbers) or unordered (with non-numbers) list

Your First R Markdown Document

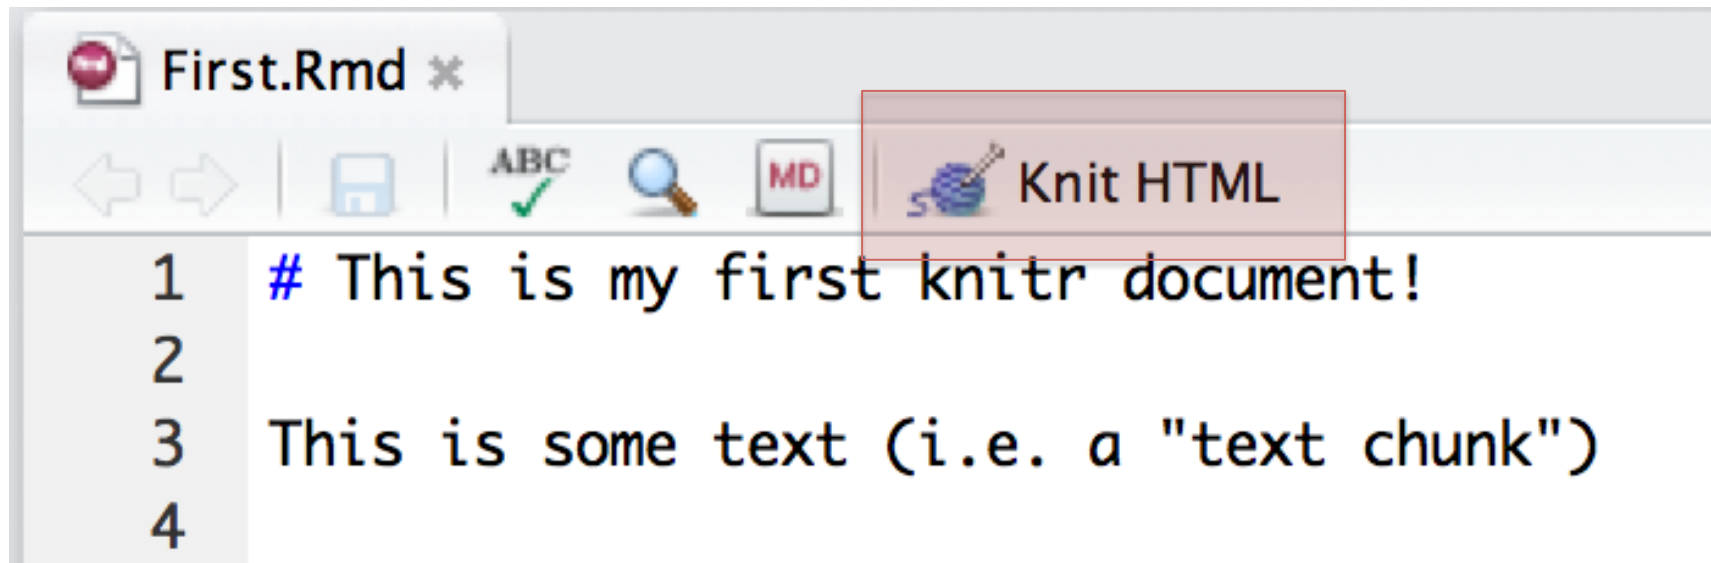
```
# This is my first knitr document!
```

This is some text (i.e. a "text chunk")

Here is a code chunk

```
```{r}
set.seed(1)
x <- rnorm(100)
mean(x)
```
```

Processing R Markdown



Resulting HTML Document

This is my first knitr document!

This is some text (i.e. a “text chunk”)

Here is a code chunk

```
set.seed(1)
x <- rnorm(100)
mean(x)
```

```
## [1] 0.1089
```

Under the Hood

```
<!-- R syntax highlighter -->
<script type="text/javascript">
var hljs=new function(){function m(p){return p.replace(/&/gm,"&
hljs.initHighlightingOnLoad();
</script>
```

```
</head>
```

```
<body>
```

```
<h1>This is my first knitr document!</h1>
```

```
<p>This is some text (i.e. a &ldquo;text chunk&rdquo;)</p>
```

```
<p>Here is a code chunk</p>
```

```
<pre><code class="r">set.seed(1)
x &lt;- rnorm(100)
mean(x)
</code></pre>
```

```
<pre><code>## [1] 0.1089
</code></pre>
```

```
</body>
```


A Few Notes

- Code chunks begin with ```{r}` on a line by itself and end with ```` on a line by itself
 - All R code goes in between
- Code chunks can have names, which is useful when making plots
- By default, code in a code chunk will be **echoed**, as will the **result** of the computation (if there is anything to print)

Processing R Markdown Documents

- The order is
 - First.Rmd (written by you!)
 - First.md (created by knitr)
 - First.html (created by knitr)
- The .md file (and the .html file) are not things we care about and **should not be edited**
- Always edit the .Rmd file

Chunk Options

```
# This is my first knitr document!
```

This is some text (i.e. a "text chunk")

Here is a code chunk

```
```{r,echo=FALSE}  
set.seed(1)
x <- rnorm(100)
mean(x)
```
```

Chunk Options

```
# This is my first knitr document!
```

This is some text (i.e. a "text chunk")

Here is a code chunk (but it doesn't print anything!)

```
```{r,echo=FALSE,results="hide"}  
set.seed(1)
x <- rnorm(100)
mean(x)
```
```

Chunk Options

This is my first knitr document!

This is some text (i.e. a “text chunk”)

Here is a code chunk (but it doesn't print anything!)

Inline Computation

```
# This is my first knitr document!
```

This is some text (i.e. a "text chunk")

Here is a code chunk (but it doesn't print anything!)

```
```{r computetime,echo=FALSE}  
time <- format(Sys.time(), "%a %b %d %X %Y")
rand <- rnorm(1)
```
```

The current time is `r time`. My favorite random number is `r rand`.

(NOTE: Look at `?strptime` for a list of all the date/time format codes.)

Inline Computation

This is my first knitr document!

This is some text (i.e. a “text chunk”)

Here is a code chunk (but it doesn't print anything!)

The current time is Thu Apr 25 11:35:42 2013. My favorite random number is 1.0894.

Graphics

```
# This is my first knitr document!
```

This is some text (i.e. a "text chunk")

Here is a code chunk

```
```{r computestuff}  
x <- rnorm(100)
y <- x + rnorm(100, sd = 0.5)
```
```

Here is a scatterplot of my data.

```
```{r scatterplot,fig.width=8, fig.height=4}  
par(mar = c(5, 4, 1, 1), las = 1)
plot(x, y, main = "My Data")
```
```


Graphics

This is my first knitr document!


This is some text (i.e. a "text chunk")

Here is a code chunk

```
x <- rnorm(100)
y <- x + rnorm(100, sd = 0.5)
```

Here is a scatterplot of my data.

```
par(mar = c(5, 4, 1, 1), las = 1)
plot(x, y, main = "My Data")
```



Graphics: What knitr Produces

```
<body>
<h1>This is my first knitr document!</h1>

<p>This is some text (i.e. a &ldquo;text chunk&rdquo;)</p>

<p>Here is a code chunk (but it doesn't print anything!)</p>

<pre><code class="r">x &lt;- rnorm(100)
y &lt;- x + rnorm(100, sd = 0.5)
</code></pre>

<p>Here is a scatterplot of my data.</p>

<pre><code class="r">par(mar = c(5, 4, 1, 1), las = 1)
plot(x, y, main = &quot;My Data&quot;);
</code></pre>

<p></p>

</body>

</html>
```

Setting Global Options

```
# This is my first knitr document!
```

First, let's show some global options.

```
```${r}  
opts_chunk$set(echo = FALSE)
```
```

This is some text (i.e. a "text chunk")

Here is a code chunk

```
```${r computestuff, echo=TRUE}  
x <- rnorm(100)
y <- x + rnorm(100, sd = 0.5)
```
```

Here is a scatterplot of my data.

```
```${r scatterplot,fig.width=8, fig.height=4}  
par(mar = c(5, 4, 1, 1), las = 1)
plot(x, y, main = "My Data")
```
```

Setting Global Options

This is my first knitr document!

First, let's show some global options.


```
opts_chunk$set(echo = FALSE)
```

This is some text (i.e. a "text chunk")

Here is a code chunk

```
x <- rnorm(100)
y <- x + rnorm(100, sd = 0.5)
```

Here is a scatterplot of my data.



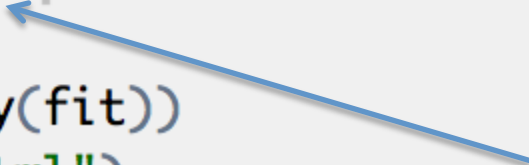
Making Tables with xtable

```
# This is my first knitr document!
```

```
```{r fitmodel}  
library(datasets)
data(airquality)
fit <- lm(Ozone ~ Wind + Temp + Solar.R, data = airquality)
```
```

Here is a table of regression coefficients.

```
```{r, results="asis"}  
library(xtable)
xt <- xtable(summary(fit))
print(xt, type = "html")
```
```



Very important!

Making Tables with xtable

This is my first knitr document!

```
library(datasets)
data(airquality)
fit <- lm(Ozone ~ Wind + Temp + Solar.R, data = airquality)
```

Here is a table of regression coefficients.

```
library(xtable)
xt <- xtable(summary(fit))
print(xt, type = "html")
```

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|-----------------|-------------------|----------------|---------------------|
| (Intercept) | -64.3421 | 23.0547 | -2.79 | 0.0062 |
| Wind | -3.3336 | 0.6544 | -5.09 | 0.0000 |
| Temp | 1.6521 | 0.2535 | 6.52 | 0.0000 |
| Solar.R | 0.0598 | 0.0232 | 2.58 | 0.0112 |

Caching Long Computations

- Important for long documents with complex computations
- Set 'cache = TRUE' in chunk option to cache results
- Code has to be run once; results are then stored in a key-value database
- On subsequent runs, results are loaded from database rather than execute code
- If code changes, results are re-run
- Dependencies can be specified via 'dependson'

Caching Long Computations

```
# This is my first knitr document!
```

```
```{r fitmodel,cache=TRUE}  
library(datasets)
data(airquality)
fit <- lm(Ozone ~ Wind + Temp + Solar.R, data = airquality)
Sys.sleep(5)
```
```

Here is a table of regression coefficients.

```
```{r, results="asis"}  
library(xtable)
xt <- xtable(summary(fit))
print(xt, type = "html")
```
```


Caching Long Computations

```
# This is my first knitr document!
```

```
```{r fitmodel,cache=TRUE}  
library(datasets)
data(airquality)
print("hello")
fit <- lm(Ozone ~ Wind + Temp + Solar.R, data = airquality)
Sys.sleep(5)
```
```

Here is a table of regression coefficients.

```
```{r, results="asis",cache=TRUE,dependson="fitmodel"}  
library(xtable)
xt <- xtable(summary(fit))
print(xt, type = "html")
Sys.sleep(2)
```
```

Summary of Basic Options

- Controlling Output
 - results: “markup”, “asis”, “hide”
 - echo: TRUE, FALSE
 - eval: TRUE, FALSE
- Figures
 - fig.width: width of plot (passed to graphics dev)
 - fig.height: height of plot
 - fig.path: directory to put figures (“figure/”)

Reproducibility: The Bigger Picture

- Reproducibility is not just about using individual tools
- Need a “reproducible workflow” that incorporates a number of other tools/practices
- Version control: git, svn, cvs (!)
- Software versioning, regression/unit testing

Thank you!